



The governed intelligence layer for business AI.

One correct answer, from your own systems. Inside ChatGPT, Claude, Perplexity, Slack, and Teams.

Seed Round 2026

315

users

148

MCP connections

40+

business connectors

PROBLEM

Two assistants. Same business. Different answers.

ChatGPT says revenue is up. Claude says it is flat. Same systems, same period. Nobody trusts the number.

ChatGPT

Revenue is up

QuickBooks, Shopify, GA4

Same business data

Different answers

Claude

Revenue is flat

QuickBooks, Shopify, GA4

Investor read. The pain is not access. The pain is correctness.

The model decides too much.

AI clients choose how business data is retrieved, interpreted, summarized, and presented. That is the missing control layer.

01

Tool choice

Which connector or tool gets called, and in what order.

02

Interpretation

Which metric definition gets used when two systems disagree.

03


Presentation

Which answer format the user finally sees and acts on.

There is no governed layer between business data and the AI.

Ask your business a question. Get one correct answer.


CorpusIQ is the governed intelligence layer inside ChatGPT, Claude, and Perplexity. The user still talks to the AI. CorpusIQ owns the answer chain.



Every answer resolves a versioned, auditable workflow. Not a fresh improvisation each time.

One question becomes an inspectable operating answer.

The verified answer renders inside the AI interface, with source-level drill-down.



“How is my business doing this month?”

Verified

7 systems checked, 0 conflicts

Pinned

Executive Snapshot v4.2.0

Inspectable

Click any card to drill into the source

Start with operators who need one trusted answer.

The beachhead is not every AI user. It is businesses already running on multiple systems and already using AI to make decisions.

● Initial customer

SMB and mid-market operators running 5+ business systems, using ChatGPT or Claude as a work interface.

● Initial buyer

CEO, COO, finance, revenue operations, ecommerce and growth leaders.

● Expansion

Regulated and enterprise workflows where answers must be reproducible, sourced, and auditable.

6.4 million US small businesses have employees, and multi-system operation is now the default. Source: SBA, 2026.

REVENUE AT STATED PENETRATION. NOT TOP-DOWN MARKET SIZE.

\$14.4M

Beachhead opportunity

0.16% of base. 10K x \$120 ARPA x 12 mo

\$300M

Expansion, SMB + mid-market

100K accounts x \$250 ARPA x 12 mo

\$60M

Expansion, enterprise

5K accounts x \$12K ACV

Expansion logic. Individual operator, to team workspace, to company-wide answer layer, to enterprise governance.

Usage is moving from connection to governed execution.

Pre-revenue. Real users, real MCP activations, real business systems connected.

315

users

148

MCP connections

47%

connection rate

Efficiency, secondary: 126.7M tokens saved through server-side aggregation instead of raw data returned to the model.

● Organic demand

Google organic signups are growing month over month.

● Channel proof

Live on the ChatGPT Apps directory, the Anthropic connector directory, and Azure Marketplace.

● Next proof point

Repeat usage, retention, and first paid conversion.

Seed signal. The next constraint is activation, training, and paid conversion.

Live where the buyer already works.

Three assistants, two team surfaces, and an enterprise procurement channel. All shipping today.

● AI assistants

ChatGPT, Claude, and Perplexity. Listed in the Anthropic connector directory and live on the ChatGPT Apps directory.

● Team surfaces

Slack and Microsoft Teams, live. The governed answer reaches the thread where the decision is actually made.

● Cloud marketplace

Azure Marketplace, approved and live. Buyable against committed Azure spend.

● Next channels

AWS and Google Cloud marketplaces through the signed Tackle.io distribution agreement.

Azure Marketplace turns enterprise procurement into a purchase order against budget already committed.

Why it compounds. Every surface is the same governed answer chain, not a separate integration. One contract, five places to ask the question, one number that comes back.

Connectors are table stakes. Governance is the moat.

The structural differentiator is reducing the AI from many execution decisions to one model-facing entry point.

● One model-facing tool

The model makes one decision. CorpusIQ resolves the governed workflow by ID and version.

● Reconciliation

Every number is cross-checked. Conflicts are surfaced. Metric definitions stay pinned.

● Fail-closed correctness

Canonical metric, validation, analysis, then visualization. Never raw data to LLM to dashboard.

● Interactive rendering

The verified answer renders inside the host, with source-level drill-down.

WHERE OTHERS STOP

● Model vendors

Own the assistant interface and first-party controls.

● MCP aggregators

Own connectivity and auth plumbing.

● BI + semantic layers

Own dashboards and structured metrics.

● CorpusIQ

Owns the governed path from question to verified answer.

Platform risk, stated plainly. Hosts may ship workflow controls. Model neutrality is shipped, not claimed: one governed answer chain across five surfaces. CorpusIQ owns what no host builds per tenant, metric definitions, reconciliation logic, governed runbooks, and audit evidence.

Operators who have built, scaled, and exited.

Seed capital is directed at conversion, onboarding, reliability, and audit-ready execution.

TEAM



Benoit Pecqueur

2 prior exits

Founder + CEO

AdDynamix, scaled to \$24M revenue and 78 people, acquired by Ybrant Digital.
WakeBallast, acquired 2023, supplied BAE Systems, Boeing, and NASA. Operator across SaaS, manufacturing, ecommerce, and real estate.



Ted Malone

15 years at Microsoft

CTO

Enterprise architecture, MCP infrastructure, OAuth, and connectors.

USE OF FUNDS

GTM + conversion

45%

Customer onboarding + training

25%

Platform reliability + security

10%

Auditable regulated workflows

10%

Compliance + operations

10%

18-month target. Repeatable paid conversion, stronger retention, enterprise-ready audit evidence, and proof for the next institutional round.

Seed \$1.0M.

SAFE. \$6.0M post-money cap. Build the governed intelligence layer for the agent economy.

● GTM conversion

Turn live adoption into repeatable paid accounts.

● Customer training

Make activation and governed workflows easy to adopt.

● Auditable answers

Source-backed execution for regulated workflows.

● 18-month proof

Conversion, retention, audit evidence, next-round readiness.

AI owns the conversation.

CorpusIQ owns the governed business answer.

Pushing governed execution into the MCP conversation.

runbook.v1 is an open governance contract for deterministic, versioned, auditable workflow execution.

● Open contract

runbook.v1 published as an application-layer contract for MCP hosts.

● Internal conformance

24 of 24 conformance tests passing against the current specification. Self-administered.

● Community proposal

Proposal submitted to the MCP repository as issue 3270. Open, not adopted.

Early signal, not standard status.

Why it matters

The intent is to make governed execution a portable contract across AI hosts, not a CorpusIQ-only behavior. Same playbook as the protocol layer itself: open contract, conformance, community implementation.